# Micro-Optimizing Go Code

•••

George Tankersley

@gtank

Code: https://github.com/gtank/blake2s

## This is a story of getting a little carried away

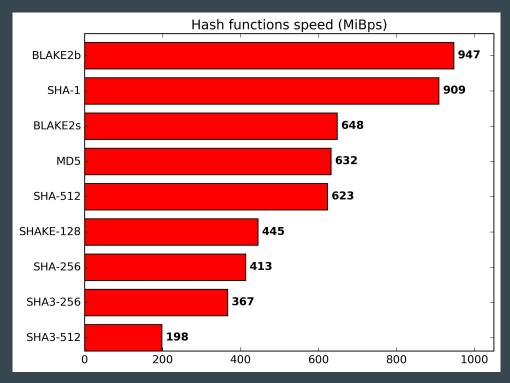
name Hash8Bytes-4 Hash1K-4 Hash8K-4	old time/op 971ns ± 4% 10.2μs ±11% 77.0μs ± 4%	new time/op 392ns ± 1% 3.1μs ± 3% 23.4μs ± 1%	delta -59.66% -69.26% -69.60%	(p=0.008) (p=0.008) (p=0.008)
name Hash8Bytes-4 Hash1K-4	old speed 8.24MB/s ± 4% 101MB/s ±10%	new speed 20.41MB/s ± 1% 327MB/s ± 3%	delta +147.65% +224.31%	(p=0.008) (p=0.008)
Hash8K-4	101MB/S ±10% 106MB/S ± 4%	$350MB/s \pm 3\%$		(p=0.008)

# BLAKE2

#### BLAKE2 is awesome

#### From the paper:

- Faster than MD5
- Immune to length extension attacks
- FEATURES! Parallelism, tree hashing, prefix-MAC, personalization, etc



Single-core serial implementation, Skylake

## BLAKE2 is *under-specified*

No one implements all of it. Not even <a href="RFC7693"><u>RFC7693</u></a>:

Note: [The BLAKE2 paper] defines additional variants of BLAKE2 with features such as salting, personalized hashes, and tree hashing. These OPTIONAL features use fields in the parameter block that are not defined in this document.

Two cryptographers implementing an unspecified algorithm.

Photo by Zach Weinersmith, circa 2009.



#### The BLAKE2 Algorithm (Abridged)

- 1. Initialize parameters
- 2. Split input data into fixed-size blocks
- 3. Scramble the bits around
- 4. Update internal state
- 5. Finalize & output

#### Hash functions in Go

```
type Hash interface {
    // Write (via the embedded io.Writer) adds more data to the hash.
    // It never returns an error.
    io.Writer
    // Sum appends the current hash to b and returns the resulting slice.
    // It does not change the underlying hash state.
    Sum(b []byte) []byte
    // Reset resets the Hash to its initial state.
    Reset()
    // Size returns the number of bytes Sum will return.
    Size() <u>int</u>
    // BlockSize returns the hash's underlying block size.
    // The Write method must be able to accept any amount
    // of data, but it may operate more efficiently if all writes
    // are a multiple of the block size.
    BlockSize() <u>int</u>
```

#### Hash functions in Go

Sum(b []byte) []byte

// It never returns an error.

// are a multiple of the block size.

type Hash interface {

io.Writer

Reset()

Size() <u>int</u> ←

```
// Write (via the embedded io.Writer) adds more data to the hash.
                                                   Mutating finalize()
       // Sum appends the current hash to be and returns the resulting slice.
leeds key It does not change the underlying hash state.
       Reset resets the Hash to its initial state.
       // Size returns the number of bytes Sum will return.
       // BlockSize returns the hash's underlying block size.
       // The Write method must be able to accept any amount
```

BlockSize() int Differs by BLAKE2 variant

// of data, but it may operate more efficiently if all writes

# Benchmarking

#### Tools of the trade

go bench <a href="https://dave.cheney.net/2013/06/30/how-to-write-benchmarks-in-go">https://dave.cheney.net/2013/06/30/how-to-write-benchmarks-in-go</a>

benchstat <a href="https://godoc.org/golang.org/x/perf/cmd/benchstat">https://godoc.org/golang.org/x/perf/cmd/benchstat</a>

pprof <a href="https://golang.org/pkg/runtime/pprof/">https://golang.org/pkg/runtime/pprof/</a>

#### Tools of the trade

go bench <a href="https://dave.cheney.net/2013/06/30/how-to-write-benchmarks-in-go">https://dave.cheney.net/2013/06/30/how-to-write-benchmarks-in-go</a>

benchstat <a href="https://godoc.org/golang.org/x/perf/cmd/benchstat">https://godoc.org/golang.org/x/perf/cmd/benchstat</a>

pprof <a href="https://golang.org/pkg/runtime/pprof/">https://golang.org/pkg/runtime/pprof/</a>

And this awful bash one-liner:

```
DATE=`date -u +'%s' | tr -d '\n'`; BRANCH=`git rev-parse --abbrev-ref HEAD`; for i in {1..8}; do go test -bench . >> benchmark-$BRANCH-$DATE; done
```

#### **Benchmarks**

- Go has built-in support for benchmarking.
- You've seen testing.T, this is <u>testing.B</u>.
- I usually put benchmarks in my test files.

The benchmarks I'm using are here:

https://github.com/gtank/blake2s/blob/master/blake2s\_test.go

#### **Benchmarks**

```
var emptyBuf = make([]byte, 8192)
func benchmarkHashSize(b *testing.B, size int) {
    b.SetBytes(int64(size))
    sum := make([]byte, 32)
    b.ResetTimer()
    for i := 0; i < b.N; i++ {
        digest, := NewDigest(nil, nil, nil, 32)
        digest.Write(emptyBuf[:size])
        digest.Sum(sum[:0])
func BenchmarkHash8Bytes(b *testing.B) {
    benchmarkHashSize(b, 8)
```

#### **Benchmarks**

```
var emptyBuf = make([]byte, 8192)
func benchmarkHashSize(b *testing.B, size int) {
    b.SetBytes(int64(size))
    sum := make([]byte, 32)
                                  MAGIC
    b.ResetTimer()
   for i := 0; i < b.N; 1++ {
        digest, := NewDigest(nil, nil, nil, 32)
        digest.Write(emptyBuf[:size])
        digest.Sum(sum[:0])
func BenchmarkHash8Bytes(b *testing.B) {
    benchmarkHashSize(b, 8)
```

#### \$ go test -bench .

```
goos: linux
goarch: amd64
pkg: github.com/gtank/blake2
BenchmarkHash8Bytes-4
                         2000000
                                     859 ns/op
                                                  9.31 MB/s
                                    8822 ns/op
BenchmarkHash1K-4
                          200000
                                                116.06 MB/s
BenchmarkHash8K-4
                                   66617 ns/op
                                                122.97 MB/s
                           20000
PASS
ok
      github.com/gtank/blake2 6.613s
```

#### \$ go test -bench .

```
goos: linux
goarch: amd64
pkg: github.com/gtank/blake2
                                    859 ns/op
BenchmarkHash8Bytes-4
                        2000000
BenchmarkHash1K-4
                                   8822 ns/op 116.06 MB/s
                         200000
BenchmarkHash8K-4
                                  66617 ns/op
                          20000
PASS
ok
      github.com/gtank/blake2 6.613s
```

# pprof

## (pprof) top5

```
Showing top 5 nodes out of 39
 flat flat%
               sum%
                             cum%
                       cum
                                   blake2.(*Digest).compress
3480ms 54.12% 54.12% 5130ms 79.78%
1320ms 20.53% 74.65% 1600ms 24.88%
                                   github.com/gtank/blake2.g
      4.35% 79.00% 280ms
                           4.35%
                                   math/bits.RotateLeft32
 280ms
 220ms 3.42% 82.43% 780ms 12.13%
                                   runtime.mallocgc
100ms 1.56% 83.98% 650ms 10.11%
                                   runtime.makeslice
```

What's their relationship though?

#### The round function, g()

```
func g(a, b, c, d, m0, m1 uint32) (uint32, uint32, uint32,
uint32) {
   a = a + b + m0
   d = bits.RotateLeft32(d^a, -16)
   c = c + d
   b = bits.RotateLeft32(b^c, -12)
   a = a + b + m1
   d = bits.RotateLeft32(d^a, -8)
   c = c + d
   b = bits.RotateLeft32(b^c, -7)
   return a, b, c, d
```

Inlining is copying the body of a function into the body of the caller.

Avoids function call overhead, which is substantial in Go.

Tradeoff between performance and binary size.

The *inliner* is a component of the <u>compiler</u> with no\* manual control.

It uses an AST visitor to calculate a complexity score vs a complexity budget.

Chasing the inliner is a flavor of optimization unique to Go.

<sup>\*</sup>Except unofficial pragmas

Functions accrue +1 cost for each node in the instruction tree

Slices are expensive! A slice node is +2 or +3 depending.

Function calls OK in most cases if we have budget for it. But a call is +2 regardless.

Some things are hard stops:

- Nonlinear control flow for, range, select, break, defer, type switch
- Recover (but not panic)
- Certain runtime funcs and all non-intrinsic assembly [#17373]

Full details (as of go1.11) in inl.go

#### Results: \$ benchstat baseline inlinable\_g

```
old time/opnew time/opdelta
                                              (p=0.000)
Hash8B-4 772ns \pm 2% 574ns \pm 0%
                                     -25.71%
Hash1K-4 8.50\mus ± 3% 5.20\mus ± 2% -38.80%
                                             (p=0.000)
Hash8K-4 65.8\mus ± 4% 39.3\mus ± 2%
                                    -40.25%
                                              (p=0.000)
      old speed new speed
                            delta
name
Hash8B-4 10.4MB/s \pm 2\% 13.9MB/s \pm 0\%
                                                  (p=0.000)
                                        +34.52%
Hash1K-4 121MB/s \pm 3% 197MB/s \pm 2% \pm 63.36%
                                                  (p=0.000)
                                                  (p=0.000)
Hash8K-4 125MB/s \pm 4% 209MB/s \pm 2% \pm 67.33%
```

# How do we check?

# \$ go test -gcflags="-m=2" 2>&1 | grep "too complex"

```
[\ldots]
./blake2s.go:272:6: cannot inline g: function too
complex: cost 133 exceeds budget 80
./blake2s.go:284:6: cannot inline NewDigest:
function too complex: cost 332 exceeds budget 80
./blake2s.go:340:6: cannot inline (*Digest).Sum:
function too complex: cost 100 exceeds budget 80
[\ldots]
```

## The round function, g()

```
func g(a, b, c, d, m0, m1 uint32) (uint32, uint32, uint32, uint32)
   a = a + b + m0
   d = bits.RotateLeft32(d^a, -16)
   c = c + d
   b = bits.RotateLeft32(b^c, -12)
   a = a + b + m1
   d = bits.RotateLeft32(d^a, -8)
   c = c + d
   b = bits.RotateLeft32(b^c, -7)
   return a, b, c, d
```

## The round function, g()

```
func g(a, b, c, d, m0, m1 uint32) (uint32, uint32, uint32, uint32)
    a = a + b + m0
    d = ((d ^ a) >> 16) | ((d ^ a) << (32 - 16))
    c = c + d
    b = ((b \land c) >> 12) \mid ((b \land c) << (32 - 12))
    a = a + b + m1
    d = ((d ^ a) >> 8) | ((d ^ a) << (32 - 8))
    c = c + d
    b = ((b \land c) >> 7) | ((b \land c) << (32 - 7))
    return a, b, c, d
```

# \$ go test -gcflags="-m=2" 2>&1 | grep "too complex"

```
[\ldots]
./blake2s.go:270:6: cannot inline g: function too
complex: cost 81 exceeds budget 80
./blake2s.go:282:6: cannot inline NewDigest:
function too complex: cost 332 exceeds budget 80
./blake2s.go:338:6: cannot inline (*Digest).Sum:
function too complex: cost 100 exceeds budget 80
[\ldots]
```

## The round function, g()

```
func g(a, b, c, d, m0, m1 uint32) (uint32, uint32, uint32, uint32)
    a = a + b + m0
    d = ((d ^ a) >> 16) | ((d ^ a) << (32 - 16))
    c = c + d
    b = ((b \land c) >> 12) \mid ((b \land c) << (32 - 12))
    a = a + b + m1
    d = ((d ^ a) >> 8) | ((d ^ a) << (32 - 8))
    c = c + d
    b = ((b \land c) >> 7) | ((b \land c) << (32 - 7))
    return a, b, c, d
```

#### Change the API!

```
func g(a, b, c, d, m1 uint32) (uint32, uint32, uint32, uint32)
   // a = a + b + m0
   d = ((d ^ a) >> 16) | ((d ^ a) << (32 - 16))
    c = c + d
    b = ((b \land c) >> 12) \mid ((b \land c) << (32 - 12))
    a = a + b + m1
    d = ((d ^ a) >> 8) | ((d ^ a) << (32 - 8))
    c = c + d
    b = ((b \land c) >> 7) | ((b \land c) << (32 - 7))
    return a, b, c, d
```

#### Change the API!

```
func g(a, b, c, d, m1 uint32) (uint32, uint32, uint32, uint32)
   // a = a + b + m0
   d = ((d ^ a) >> 16) ((d ^ a) << (32 - 16))
    c = c + d
    b = ((b \land c) >> 12) \mid ((b \land c) << (32 - 12))
    a = a + b + m1
    d = ((d ^ a) >> 8) | ((d ^ a) << (32 - 8))
    c = c + d
    b = ((b \land c) >> 7) \mid ((b \land c) << (32 - 7))
    return a, b, c, d
```

## \$ go test -gcflags="-m=2" 2>&1 | grep "can inline"

```
./blake2s.go:270:6: can inline g with cost 74 as:
func(uint32, uint32, uint32, uint32)
(uint32, uint32, uint32) { d = (d ^ a) >>
16 (d ^ a) << (32 - 16); c = c + d; b = (b ^ c)
\Rightarrow 12 | (b ^ c) << (32 - 12); a = a + b + m1; d =
(d ^ a) >> 8 | (d ^ a) << (32 - 8); c = c + d; b
= (b ^ c) >> 7 | (b ^ c) << (32 - 7); return a,
b, c, d }
```

# Under budget does not mean faster!

#### Don't be this guy\*

```
// This function is written to ensure it inlines.
func mul64x64(lo, hi, a, b uint64) (ol uint64, oh uint64) {
    t1 := (a>>32)*(b&0xFFFFFFFF) + ((a & 0xFFFFFFFF) * (b & 0xFFFFFFFFF) >> 32)
    t2 := (a&0xFFFFFFFFFF)*(b>>32) + (t1 & 0xFFFFFFFFF)
    ol = (a * b) + lo
    cmp := ol < lo
    oh = hi + (a>>32)*(b>>32) + t1>>32 + t2>>32 + uint64(*(*byte)(unsafe.Pointer(&cmp)))
    return
}
```

What's next?

#### Back to pprof

We need a more granular view...

#### What's going on here?

```
1s
              1s
                              m0 := m[SIGMA[round][2*i+0]]
210ms
           210ms
                   506d21:
                                    MOVO DX, CX
 70ms
            70ms
                   506d31:
                                    MOVQ DX, BX
                                    SHLQ $0x1, DX
150ms
           150ms
                   506d34:
                   506d37:
                                    CMPQ $0×10, DX
                   506d3b:
                                    JAE 0x5071cb
40ms
                   506d41:
                                    MOVQ CX, DX
            40ms
90ms
                   506d44:
                                    SHLQ $0x6, CX
            90ms
170ms
           170ms
                   506d48:
                                    LEAQ github.com/gtank/blake2s.SIGMA(SB), SI
 10ms
            10ms
                   506d4f:
                                    ADDQ SI, CX
90ms
            90ms
                   506d52:
                                    MOVL O(CX)(BX*8), DI
                                    CMPQ $0×10, DI
100ms
           100ms
                   506d55:
                   506d59:
                                    JAE 0x5071cb
                                    MOVL 0x58(SP)(DI*4), DI
 70ms
            70ms
                   506d5f:
                                         runtime.panicindex(SB)
                   5071cb:
                   5071d0:
                                    UD2
                                                  Bounds checking!
```

#### runtime.panicindex

```
$ go run bce.go
panic: runtime error: index out of range
goroutine 1 [running]:
main.demo(...)
        /home/gtank/bce.go:9 main.main()
        /home/gtank/bce.go:5 +0x11
exit status 2
```

#### Another view: \$ go test -gcflags="-d=ssa/check\_bce/debug=1"

```
\lceil \dots \rceil
./blake2s.go:199:11: Found IsInBounds
./blake2s.go:199:24: Found IsInBounds
./blake2s.go:200:11: Found IsInBounds
./blake2s.go:200:24: Found IsInBounds
197 ▶ for round := 0; round < RoundCount; round++ {
          m0 := m[SIGMA[round][2*i+0]]
     ▶  m1 := m[SIGMA[round][2*i+1]]
             switch i {
```

#### Bounds check elimination, normally

```
func (bigEndian) PutUint32(b []byte, v uint32) {
   _ = b[3] // early bounds check to guarantee safety below
   b[0] = byte(v >> 24)
   b[1] = byte(v >> 16)
   b[2] = byte(v >> 8)
   b[3] = byte(v)
```

#### Optimizing that table lookup

Combination of several "old-school" optimization techniques:

- Propagate constants
- Unroll loops
- Reuse previously-allocated local variables

In pursuit of a specific thing:

Bounds-Check Elimination (<u>further reading</u>)

```
Results - worth it?
Much, much faster!
Macros would be nice here.
```

```
v0, v4, v8, v12 = g(v0+v4+m0, v4, v8, v12, m1)
v1, v5, v9, v13 = g(v1+v5+m2, v5, v9, v13, m3)
v2, v6, v10, v14 = g(v2+v6+m4, v6, v10, v14, m5)
v3, v7, v11, v15 = g(v3+v7+m6, v7, v11, v15, m7)
v0, v5, v10, v15 = g(v0+v5+m8, v5, v10, v15, m9)
v1, v6, v11, v12 = g(v1+v6+m10, v6, v11, v12, m11)
v2, v7, v8, v13 = g(v2+v7+m12, v7, v8, v13, m13)
v3, v4, v9, v14 = g(v3+v4+m14, v4, v9, v14, m15)
v0, v4, v8, v12 = g(v0+v4+m14, v4, v8, v12, m10)
v1, v5, v9, v13 = g(v1+v5+m4, v5, v9, v13, m8)
v2, v6, v10, v14 = g(v2+v6+m9, v6, v10, v14, m15)
v3, v7, v11, v15 = g(v3+v7+m13, v7, v11, v15, m6)
v0, v5, v10, v15 = g(v0+v5+m1, v5, v10, v15, m12)
v1, v6, v11, v12 = g(v1+v6+m0, v6, v11, v12, m2)
v2, v7, v8, v13 = g(v2+v7+m11, v7, v8, v13, m7)
v3, v4, v9, v14 = g(v3+v4+m5, v4, v9, v14, m3)
v0, v4, v8, v12 = g(v0+v4+m11, v4, v8, v12, m8)
v1, v5, v9, v13 = g(v1+v5+m12, v5, v9, v13, m0)
v2, v6, v10, v14 = g(v2+v6+m5, v6, v10, v14, m2)
v3, v7, v11, v15 = g(v3+v7+m15, v7, v11, v15, m13)
```

## Results: \$ benchstat inlinable\_g eliminate\_bounds\_checks

name Hash8Bytes-4 Hash1K-4 Hash8K-4	old time/op 574ns ± 0% 5.20μs ± 2% 39.3μs ± 2%	new time/op 420ns ± 2% 2.91μs ± 3% 21.5μs ± 4%	delta -26.90% -44.09% -45.16%	(p=0.000) (p=0.000) (p=0.000)
name	old speed	new speed	delta	
Hash8Bytes-4	13.9MB/s ± 0%	$19.1MB/s \pm 2\%$	+36.81%	(p=0.000)
Hash1K-4	197MB/s ± 2%	352MB/s ± 3%	+78.87%	(p=0.000)
Hash8K-4	209MB/s ± 2%	$380MB/s \pm 4\%$	+82.40%	(p=0.000)

#### One more bounds check...?

The internal hash state is a slice, but it's always of fixed size.

Can we eliminate these?

Well, not as we expect.

#### SSA bounds check output

```
./blake2s.go:308:14: Found IsInBounds
./blake2s.go:309:14: Found IsInBounds
./blake2s.go:310:14: Found IsInBounds
./blake2s.go:311:14: Found IsInBounds
./blake2s.go:312:14: Found IsInBounds
./blake2s.go:313:14: Found IsInBounds
./blake2s.go:314:14: Found IsInBounds
```

#### Corresponding to these lines in compress():

```
307 \ d.h[0] = d.h[0] ^ v0 ^ v8

308 \ d.h[1] = d.h[1] ^ v1 ^ v9

309 \ d.h[2] = d.h[2] ^ v2 ^ v10

310 \ d.h[3] = d.h[3] ^ v3 ^ v11

311 \ d.h[4] = d.h[4] ^ v4 ^ v12

312 \ d.h[5] = d.h[5] ^ v5 ^ v13

313 \ d.h[6] = d.h[6] ^ v6 ^ v14

314 \ d.h[7] = d.h[7] ^ v7 ^ v15
```

#### Sure, why not?

Replace slice with array.

Compiler is satisfied.

No more runtime bounds checks!

Side effect: makes an explicit copy an implicit one.

```
@@ -94,7 +94,7 @@ func initFromParams(p *parameterBlock) *Digest {
       h7 := IV7 ^ binary.LittleEndian.Uint32(paramBytes[28:32])
               buf: make([]byte, 0, BlockBytes),
@@ -154,10 +154,6 @@ func (d *Digest) compress() error {
       // Split the buffer into 16x32-bit words.
       m0 := binary.LittleEndian.Uint32(d.buf[0*4 : 0*4+4])
       m1 := binary.LittleEndian.Uint32(d.buf[1*4 : 1*4+4])
@@ -328,8 +324,6 @@ func (d *Digest) finalize() ([]byte, error) {
       dCopy := *d
       dCopy.buf = make([]byte, cap(d.buf)) // want this zero-padded to BlockSize anyway
       copy(dCopy.buf, d.buf)
```

@@ -69,7 +69,7 @@ func (p \*parameterBlock) Marshal() []byte {

// The internal state of the BLAKE2s algorithm.

dCopy.t0 += uint32(len(d.buf))

diff --git a/blake2s.go b/blake2s.go
index 7b0ccd0..a349bcf 100644

--- a/blake2s.go +++ b/blake2s.go

type Digest struct {

## Results: \$ benchstat eliminate\_checks use\_fixed\_array

name	old time/op	new time/op	delta	
Hash8Bytes-4	420ns ± 2%	373ns ± 2%	-11.03%	(p=0.000)
Hash1K-4	2.91µs ± 3%	2.87µs ± 3%	~	(p=0.130)
Hash8K-4	21.5µs ± 4%	21.6µs ± 3%	~	(p=0.536)
	·	·		,
name	old speed	new speed	delta	
Hash8Bytes-4	19.1MB/s ± 2%	21.4MB/s ± 2%	+12.37%	(p=0.000)
Hash1K-4	352MB/s ± 3%	357MB/s ± 3%	~	(p=0.130)
Hash8K-4	$380MB/s \pm 4\%$	$379MB/s \pm 3\%$	~	(p=0.536)

#### Results: \$ benchstat eliminate\_checks use\_fixed\_array

name Hash8Bytes-4 Hash1K-4	old time/op 420ns ± 2% 2.91µs ± 3%	new time/op 373ns ± 2% 2.87µs ± 3%	delta -11.03% ~	(p=0.000) (p=0.130)
		•	~	<b>*</b> 1
Hash8K-4	21.5μs ± 4%	21.6μs ± 3%		(p=0.536)
name	old speed	new speed	delta	
Hash8Bytes-4	$19.1MB/s \pm 2\%$	$21.4MB/s \pm 2\%$	+12.37%	(p=0.000)
Hash1K-4	$352MB/s \pm 3\%$	$357MB/s \pm 3\%$	~	(p=0.130)
Hash8K-4	$380MB/s \pm 4\%$	$379MB/s \pm 3\%$	~	(p=0.536)



#### Benchmark

```
var emptyBuf = make([]byte, 8192)
func benchmarkHashSize(b *testing.B, size int) {
    b.SetBytes(int64(size))
    sum := make([]byte, 32)
    b.ResetTimer()
    for i := 0; i < b.N; i++ {
       digest, := NewDigest(nil, nil, nil, 32)
       digest.Write(emptyBuf[:size])
       digest.Sum(sum[:0])
```

finalize() runs once each time you calculate a BLAKE2 sum.

We eliminated a make/copy there.

```
(pprof) top10
Showing nodes accounting for 5.44s, 87.46% of 6.22s total
Dropped 61 nodes (cum <= 0.03s)
Showing top 10 nodes out of 45
 flat flat% sum%cum
                       cum%
3.46s 55.63% 55.63% 3.46s 55.63%
                                   github.com/gtank/blake2s.g (inline)
                                   github.com/gtank/blake2s.(*Digest).compress
0.82s 13.18% 68.81% 4.33s 69.61%
0.37s 5.95% 74.76% 0.96s 15.43%
                                   runtime.mallocgc
0.17s 2.73% 77.49% 0.17s 2.73%
                                   runtime.nextFreeFast (inline)
                                   github.com/gtank/blake2s.(*Digest).Write
0.16s 2.57% 80.06% 3.48s 55.95%
0.12s 1.93% 81.99% 1.42s 22.83%
                                   github.com/gtank/blake2s.(*Digest).finalize
0.11s 1.77% 83.76% 0.83s 13.34%
                                   runtime.makeslice
0.10s 1.61% 85.37% 0.10s 1.61%
                                   runtime.memmove
                                   github.com/gtank/blake2s.(*parameterBlock).Marsha
0.07s 1.13% 86.50%
                     0.26s 4.18%
                                   encoding/binary.littleEndian.Uint32 (inline)
0.06s
       0.96% 87.46%
                     0.06s
                            0.96%
```

```
$ ag "make\(" blake2s.go
55: buf := make([]byte, 32)
98: buf: make([]byte, 0, BlockSize),
325:dCopy.buf = make([]byte, cap(d.buf)) // want zero-padded to BlockSize anyway
343:out := make([]byte, dCopy.size)
390:params.Salt = make([]byte, SaltLength)
399:params.Personalization = make([]byte, SeparatorLength)
414:keyBuf := make([]byte, BlockSize)
```

```
$ ag "make\(" blake2s.go
55: buf := make([]byte, 32)
98: buf: make([]byte, 0, BlockSize),
325:dCopy.buf = make([]byte, cap(d.buf)) // want zero-padded to BlockSize anyway
343:out := make([]byte, dCopy.size)
390:params.Salt = make([]byte, SaltLength)
399:params.Personalization = make([]byte, SeparatorLength)
414:keyBuf := make([]byte, BlockSize)
```

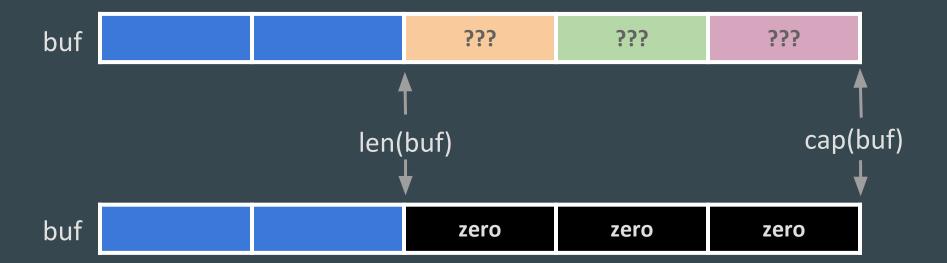
## Copy **★**

#### Struct —

```
func (d *Digest) finalize() ([]byte, error)
    if d.f0 != 0 {
        return nil, errors.New("blake2s: tr
   // make copies of everything
    dCopy := *d
    dCopy.buf = make([]byte, cap(d.buf)) //
    copy(dCopy.buf, d.buf)
```

```
type Digest struct {
    h [8]uint32
   t0, t1 uint32
    f0, f1 uint32
    buf []byte
    // size is defint
   // return. Since
    size int
```

## Zeroing the buffer



## Pattern matching (link)

```
// Zero the unused portion of the buffer. This
triggers a specific optimization for memset, see
https://codereview.appspot.com/137880043
padBuf := d.buf[len(d.buf):cap(d.buf)]
for i := range padBuf {
   padBuf[i] = 0
dCopy.buf = d.buf[0:cap(d.buf)]
```

#### Results: \$ benchstat use\_fixed\_array use\_memset

```
old time/op
                             new time/op
                                            delta
name
                 627ns ± 0%
Hash8Bytes-4
                                596ns ± 0% -4.94%
                                                    (p=0.000)
Hash1K-4
               4.28 \mu s \pm 0\% 4.24 \mu s \pm 0\% -0.90\% (p=0.000)
                31.4 \mu s \pm 0\% 31.4 \mu s \pm 0\% -0.13\% (p=0.000)
Hash8K-4
              old speed
                             new speed
                                            delta
name
Hash8Bytes-4
              12.8MB/s \pm 0\%
                            13.4MB/s \pm 0\% +5.24\%
                                                    (p=0.000)
Hash1K-4
              239MB/s \pm 0\% 241MB/s \pm 0\% +0.92\% (p=0.000)
              261MB/s \pm 0\% 261MB/s \pm 0\% +0.13\% (p=0.000)
Hash8K-4
```

```
$ ag "make\(" blake2s.go
55: buf := make([]byte, 32)
98: buf: make([]byte, 0, BlockSize),
325:dCopy.buf = make([]byte, cap(d.buf)) // want zero-padded to BlockSize
343:out := make([]byte, dCopy.size)
390:params.Salt = make([]byte, SaltLength)
399:params.Personalization = make([]byte, SeparatorLength)
414:keyBuf := make([]byte, BlockSize)
```

#### Reuse the slice

It's just slice reuse in an append API (like Sum)

```
// Sum appends the current hash to b and returns
// It does not change the underlying hash state.
func (d *Digest) Sum(b []byte) (out []byte) {
 // If there's capacity, reuse the b slice
    if n := len(b) + d.size; cap(b) >= n {
    ▶ out = b[:n]
    } else {
    out = make([]byte, n)
      copy(out, b)
    err := d.finalize(out[len(b):])
    if err != nil {
       return out[:len(b)]
    return out
```

## sliceForAppend (link)

```
// sliceForAppend takes a slice and a requested number of bytes. It returns a
// slice with the contents of the given slice followed by that many bytes and a
// second slice that aliases into it and contains only the extra bytes. If the
// original slice has sufficient capacity then no allocation is performed.
func sliceForAppend(in []byte, n int) (head, tail []byte) {
   if total := len(in) + n; cap(in) >= total {
       head = in[:total]
   } else {
       head = make([]byte, total)
       copy(head, in)
   tail = head[len(in):]
   return
```

## Results: \$ benchstat use\_memset reuse\_slices

name	old time/op	new time/op	delta	
Hash8Bytes-4	310ns ± 1%	284ns ± 0%	-8.56%	(p=0.001)
Hash1K-4	$2.73 \mu s \pm 2\%$	2.69μs ± 2%	-1.60%	(p=0.035)
Hash8K-4	$20.7\mu s \pm 3\%$	20.5μs ± 2%	~	(p=0.234)
name	old speed	new speed	delta	
Hash8Bytes-4	$25.7MB/s \pm 1\%$	28.1MB/s ± 0%	+9.36%	(p=0.001)
Hash1K-4	$375MB/s \pm 2\%$	$381MB/s \pm 2\%$	+1.63%	(p=0.038)
Hash8K-4	395MB/s ± 3%	399MB/s ± 2%	~	(p=0.234)

#### Results: \$ benchstat use\_memset reuse\_slices

```
old time/op
                             new time/op
                                            delta
name
                 310ns ± 1%
                                284ns ± 0% -8.56%
Hash8Bytes-4
                                                     (p=0.001)
                2.73\mus ± 2% 2.69\mus ± 2% -1.60%
Hash1K-4
                                                     (p=0.035)
                                                     (p=0.234)
Hash8K-4
                20.7\mu s \pm 3\%
                            20.5μs ± 2%
              old speed
                             new speed
                                            delta
name
Hash8Bytes-4
              25.7MB/s \pm 1\% 28.1MB/s \pm 0\% +9.36\%
                                                     (p=0.001)
Hash1K-4
               375MB/s \pm 2\%
                              381MB/s \pm 2\% +1.63\%
                                                     (p=0.038)
                                                     (p=0.234)
Hash8K-4
               395MB/s \pm 3\%
                              399MB/s \pm 2\%
```

## Diminishing Returns

#### Diminishing Returns

#### These look like:

- Don't allocate some trivial intermediate variables
- Unroll remaining fixed loops
- Copy small functions into this package to allow inlining them
- Hunt down the less significant bounds checks

#### Worth it? Not really.

Many hours of my life.

Library of techniques, not always best practices.

Extremely compiler version dependent.

Still not competitive with assembly.

# STRIKE FORCE

# Micro-Optimizing Go Code

•••

George Tankersley

@gtank

Code: https://github.com/gtank/blake2s